

Analyzing contingent interactions in R with `chattr`

Anonymous CogSci submission

Abstract

We present `chattr`, an R package that enables users to easily detect and describe temporal contingencies in pre-annotated interactional data. The need to detect and analyze temporal contingencies is ubiquitous across studies of signal exchange systems, including human and non-human animal communication systems. At present, approaches to this kind of analysis require careful manual evaluation (i.e., do not scale up), are proprietary and/or over-specialized (i.e., are limited in use), or are constructed ad-hoc for each study (i.e., are variable in construct). `Chattr`'s theoretically motivated, customizable, and open source code provides a set of core functions that allow users to quickly and automatically extract contingencies in data that have already been annotated for interactant activity (via manual or automated application) in one of multiple common data formats (.txt, .rttms, .its). Particularly given the rise of powerful, long-lasting recording devices and automated tools for diarizing activity in these recordings, `chattr` stands to facilitate valid at-scale analysis of temporal contingencies in a wide variety of data types while also providing consistency across studies in the analytical approach. Detailed output includes information about related turns ('prompts' and 'responses') and the allocation of turns into interactional sequences ('conversational bouts') and solo sequences ('monologues') for every emission produced by a target interactant. We here demonstrate the use of `chattr` by testing predictions about turn-taking behavior in three language development corpora, showing that the package effectively recovers documented variation in linguistic input given both manual and automatically created speech annotations. We note future directions for package development, with the hope that it will be used for the convenient management and analysis of interactional data across multiple distinct research contexts.

Keywords: Contingency; interaction; turn taking; LENA; communication; conversation; R; automated speech annotation.

Introduction

The current paper introduces an R (REFS) package called `chattr` that facilitates the detection and analysis of temporally contingent interactions in pre-annotated data (URL redacted).¹ The utility of this package extends across studies of adult and child human interactions, non-human animal communication, and contingencies within multi-modal signals. Despite significant common conceptual ground between these domains, definitions of contingency phenomena and implementations of contingency detection remain fairly

inconsistent, foregoing critical opportunities for the accumulation of shared construct validity across domains. In part, these divergences are due to the lack of flexible tools for contingency analysis; existing systems are either constructed ad-hoc or are proprietary, and thereby limited in who can use them. The `chattr` package aims to improve this situation in two ways: (1) It takes inspiration from the fields of conversation analysis, psycholinguistics, and language development to provide theoretically sound, but customizable measurements of temporally contingent interaction at scale and (2) its open-source toolkit accepts a handful of generic formats as input, opening up its analytical framework to a wide variety of domains (e.g., child language input, multi-party adult conversation, non-human animal signal exchanges, multi-modal contingencies produced by a single organism, etc.). The remainder of this short report reviews the theoretical basis underlying the development of `chattr`, describes the core functions offered by the package, and demonstrates its use in three existing datasets.

Contingent interaction

The joint coordination of action by two or more agents usually involves temporal contingencies. Whether we are making music with others, crossing a busy intersection, or chatting with a friend, the timing of our contributions to a coordinated event is crucial to its success. In many cases, the optimal strategy for coordination involves some form of turn taking. In a typical turn-taking interaction, only one interactant makes their contribution at a time, and decisions about who contributes when can be determined flexibly (as in conversation) or in a pre-defined manner (as in a debate). This sequential structure enables interactants to adapt each contribution such that it relevantly progresses the joint activity and to initiate unplanned subsequences (e.g., repairing a misunderstanding) without breaking from moving toward the larger goal.

Turn-taking (and similar temporally contingent) interactions are essential for communication across the animal kingdom (REFS). In humans, turn-taking interactions may be the only reliable source of language universals (REFS). Traditionally, these kinds of interactional contingencies have been studied using careful inspection and analysis (both qualitative and quantitative) of manual measurements from video and audio recordings. However, recent advances in automated an-

¹At time of submission, the conversion from R scripts to R package is still underway. That said, all documentation, scripts, and tests are available at the given URL and the fully packaged version will be available before May 2021.

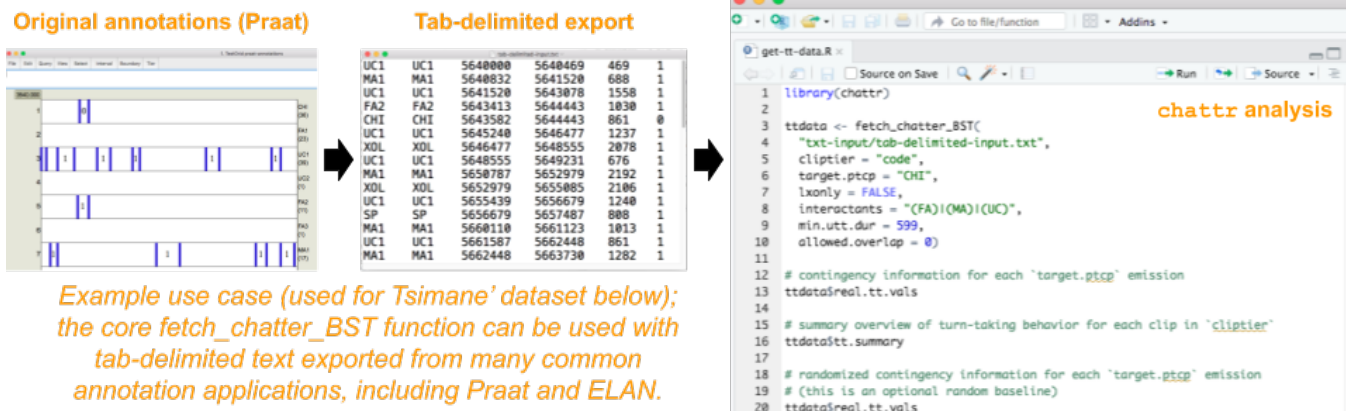


Figure 2: Example workflow for an annotation file using ‘chattr’.

Users can use the default settings for each function—which include limits on the relevant temporal windows, potential interactants, and restrictions on productions considered—or can customize as desired. More advanced users can capitalize on the wide variety of sub-functions utilized by the core input-type functions to tailor `chattr`’s functions to their unique needs. All settings, output information types, and theoretical background is thoroughly summarized in the online documentation where the project is stored on GitHub.

Core concepts We encourage users to first evaluate how well `chattr`’s concepts of ‘turn’, ‘transition’, and ‘interactional sequence’ fit those of the user; our default definitions differ from those typically used in the language development literature and are somewhat restricted compared to their full (and human conversation-specific) theoretical meanings in conversation analysis (e.g., see REFS). We briefly summarize these core concepts here. The same concepts are illustrated in Figure 1. We use the concepts of ‘producer’ and ‘recipient’/‘addressee’ rather than ‘speaker’ and ‘listener’ to underscore the utility of these concepts across modalities, species, and interactional contexts:

A ‘turn’ comprises one or more closely occurring emissions by the same producer. That is, a turn can be formed of multiple complete emissions (e.g., utterances/communicative acts) that may be separated by pauses in production so long as (a) there is no intervening emission from another producer and (b) the pause in production is short. An example of a single-unit turn in English is “Jane is the one in the hat.”. An example of a multi-unit turn in English is “Jane is the one in the hat [pause] third from the left.”.

A ‘turn transition’ occurs when one producer’s turn stops and another producer’s turn begins. Every turn transition has a pre-transition producer and a post-transition producer—these must be different individuals. The transition *begins* when the first turn ends and *ends* when the second turn starts. Therefore, if the second turn starts before the first turn ends, the transition time is negative; this is referred to as ‘transitional overlap’. If the second turn starts after the first turn

ends, the transition time is positive; referred to as ‘transitional gap’.

An ‘interactional sequence’ is an unbroken turn-taking sequence between the target interactant and one or more of their interactional partners. Interactional sequences likely index more structurally complex, engaged interactional behaviors than single turn transitions do—akin to conversational bouts (or LENA VABs)—in which participants can more substantially build on joint goals.

Default settings are designed for human spontaneous conversation, including conversation, which demonstrates fairly robust timing patterns (with some systematic variation) across the signed and spoken languages that have been analyzed (REFS). The three most critical default settings are that: (a) up to 2000 ms of transitional gap or up to 1000 ms of transitional overlap is allowed between turns, (b) turn transitions can occur with turns of any duration, content, and between any potential interactional partner, and (c) when there are multiple potential prompts or responses (e.g., two interactants answer a question nearly simultaneously), `chattr` picks the turn that begins soonest. Users interested in emulating LENA’s CTC measure with their .its files can use a specialized version of the tool in which the target producer is assumed to be “CH” (target child), potential interactants limited to “FA” and “MA” (female and male adult), and analyzed turns contain some linguistic material (REFS).

Example use case Suppose that I am interested in investigating how adult turn-taking varies in a dataset across semi-structured contexts (e.g., during board game play). I would ensure that the annotations are formatted as tab-delimited text (see the `chattr` documentation). Then I would use the Basic Speech Table call `fetch_chatter_BST()` to fetch turn-taking information. I might also desire to, e.g., employ a minimum utterance duration and a more strict temporal window for contingency, as well as calculate 10 randomized simulations of turn-taking rates to assess the baseline likelihood of finding a contingency: `fetch_chatter_BST(filename, min.utt.dur = 1500,`

allowed.gap = 1000, allowed.overlap = 600, n.runs = 10). This call yields detailed tables of detected turn-taking behavior ready for the author’s statistical analysis of choice (Figure 2).

Pilot analysis

We demonstrate the use of `chattr` with three child language environment datasets from rural Indigenous communities. These recordings document children’s verbal interactional patterns over full waking days at home in underrepresented research populations. `Chattr` allows us to examine interactional patterns at scale in these corpora, evading months of manual annotation achieving the same aim. The first two corpora, Tseltal (Mayan; Chiapas, Mexico; N = 10) and Yélf Dnye (isolate; Milne Bay, Papua New Guinea; N = 10), come from the Casillas HomeBank corpus (REFS) and were made with near parallel methods: children under age 3;0 wore an Olympus WS-832/853 audio recorder during a day at home for 8–11 hours. The third corpus, Tsimane’ (Tsimane’; Bolivia; N = 27) features children under 6;0 who wore one of multiple recording devices (LENA, Olympus, or USB) during a day at home for 4–21 hours (REFS); we focus here on the subset of those 27 recordings made with LENA (N = 17). In each dataset we assess the baseline turn-taking rate and duration of interactional sequences over age and by interactant type. For the Tsimane’ corpus we briefly compare `chattr` estimates on both LENA (automated) and manually created annotations of the same recording minutes. Here `chattr` gives a brief glimpse into how reliably known patterns from these children’s linguistic input (REFS) are recapitulated in their turn-taking behavior.

Study 1. Tseltal and Yélf Dnye

We analyze interactional behavior in 20 clips for each recording: 9 randomly selected clips (5 min for Tseltal and 2.5 min for Yélf Dnye), 5 clips manually selected for day-peak turn-taking behavior of the target child with one or more interactants (each 1 min), 5 clips manually selected for day-peak vocal activity by the target child (each 1 min), and one 5-minute expansion on the most active turn-taking or vocal activity clip. Each clip was manually annotated for all hearable speech, including addressee coding (e.g., target-child-directed vs. other-directed; see REFS for details). Despite documented differences in caregiver-child interactional style (REFS), day-long linguistic input estimates show similar directed linguistic input patterns in these two communities. While female adult speech constitutes the majority of linguistic input in both communities, Yélf children show a marked increase in directed speech from other children with age. This pattern appears more weakly in the Tseltal data. We therefore expected to find that: (1) turn-taking rates are higher in turn-taking clips than in vocal activity and random clips, (2) rates are overall similar between the two communities, and (3) interactional sequences involving other children increase with age, particularly for Yélf children.




Figure 3: Turn transition rate by corpus, divided across random (blue) and manually selected turn-taking/high-vocal-activity clips (red).




Figure 4: Proportion of interactional sequences involving at least one non-target child across age, by language.

Methods We use the `fetch_chatter_AAS()` call, which is specifically designed for those using the ACLEW Annotation Scheme (REFS)[⁴]: that allows 2000 ms of gap and 1000 ms of overlap between at turn transitions and searches over all annotated utterances (any duration, content, and from any speaker). We limit our analysis to utterances directed exclusively to the target child. We also indicate the annotated regions by using the `cliptier` argument (see documentation).

[⁴] ACLEW is a working group that has provided a flexible manual annotation template friendly for use with open

source speech processing tools (REFS).

Results The mean rate of turn transitions in the Tseltal corpus was 3 and 11.8 transitions per minute for the random and active clips, respectfully. The mean rates in the Yélf Dnye corpus were 2.4 and 12.8 for the random and active clips. Overall the distribution of turn taking rates across annotated clips was similar between the two sites (Table 1). A linear mixed effects regression of turn-transitions per minute with predictors of clip type, corpus, and their interaction and a random intercept for child reveals that, indeed, random clips have significantly lower turn-transition rates ($B = -8.78$, $SE = 1.2$, $t = -7.31$) and there is no evidence for a significant difference in turn-taking rates between languages ($t = 0.54$) and no evidence for a clip type-language interaction ($t = -0.74$).

A second linear mixed effects regression of the proportion of interactional sequences that feature at least one non-target child with predictors of age (in months), corpus, and their interaction and a random intercept for child reveals that, as expected, there is a significant age-by-corpus interaction by which Yélf children show a larger increase in other-child interactional sequences with age compared to Tseltal children ($B = 0.01$, $SE = 0.01$, $t = 2.47$). There is no evidence for simple effects of age ($t = 0.2$) or language ($t = -0.99$).

Study 2. Tsimane'

These Tsimane' recordings were first automatically analyzed with LENA and then subsequently (and independently) manually annotated in one minute clips, every 60 minutes, starting at the 34th minute (34 min, 94 min, 154 min, etc.) in Praat (REFS). Both annotation formats include information regarding (a) when speech was occurring and (b) what type of speaker produced it (i.e., the target child, a nearby woman/man/other child, or other noise sources) for each of the hand-annotated minutes. Prior analysis shows comparably low rates of directed speech in these Tsimane' data to the Tseltal and Yélf Dnye recordings, again with a high proportion of directed input coming from other children. We therefore expected to find that: (1) turn-taking rates are overall similar to what we found in the random samples of the other two communities, (2) turn-taking sequences involving other children are comparable or more frequent than found in the random samples of the other two communities, (3) interactional sequences involving other children increase with age, and (4) manual and automated speech annotations of the same audio clips result in similar turn-taking estimates.

Corpus	Clip type	mean (sd; range), median
Tseltal	active (manual)	11.8 (4.8; 4.5-20.1), 12.3
Tseltal	random (manual)	3 (3.1; 0.4-10.6), 2.3
Yélf Dnye	active (manual)	12.8 (6.5; 3.9-22.2), 10.8
Yélf Dnye	random (manual)	2.4 (1.6; 0.5-6), 2.2
Tsimane'	random (LENA)	3.2 (1.1; 1.2-5.1), 3.1
Tsimane'	random (manual)	3.2 (1.2; 1.3-6), 3




Figure 5: Turn transition rate by annotation type (LENA automated vs. manual) in the same audio clips. Clips are a periodic random sample of the daylong recording.




Figure 6: Proportion of interactional sequences involving at least one non-target child across age, by annotation type (LENA automated vs. manual) in the same audio clips.

Methods We first use the `fetch_chatter_BST()` call with the manually annotated data, matching conditions of the call as closely as possible to what can be compared in the LENA output files, that is: include woman, man, and other-child speech, both linguistic and non-linguistic, with a minimum utterance duration of 600ms (the LENA lower limit) and no overlap allowed (meaningful overlap is not possible in LENA). With the automatic LENA annotations on the same recordings (in the same 1-minute segments) we adjust the de-

faults on `fetch_chatter_LENA()` to reflect the same restrictions.

Results A linear mixed effects regression of turn-transitions per minute with predictors of annotation type (LENA automated vs. manual) and a random intercept for child reveals that turn-transition rates are similar between the two annotation methods ($B = -0.09$, $SE = 0.41$, $t = -0.23$). Consistent with our hypothesis, the mean rate of turn transitions was similar to what we found in the Tseltal and Yélf Dnye random clips, at 3.2 transitions per minute, though with fewer instances of rates above 5/min (Table 1).

A second linear mixed effects regression of the proportion of interactional sequences that feature at least one non-target child with predictors of age (in months), annotation type (LENA vs. manual), and their interaction and a random intercept for child reveals that, as expected, there is a significant increase in other-child interactional sequences with age ($B = 0.01$, $SE = 0$, $t = 3.39$). There is no evidence for simple effects of annotation type ($t = 0.71$) or for an age-annotation type interaction ($t = -1.39$).

Contribution and next steps

The `chattr` package allows users to easily implement theoretically informed contingency analyses on a wide variety of data types, including both automatically and manually annotated data. The package is designed for both straightforward (i.e., basic `fetch_chatter` calls) and customized analysis scenarios and provides detailed outputs that can be merged with other data about the same recordings. By providing a single tool for analyzing the most common input formats used for interactional data in psychology, animal behavior, and speech technology research, `chattr` aims to help build theoretical and methodological connections regarding the nature of contingent behaviors across diverse domains. While `chattr` has now been tested on a variety of child language datasets, new functionality will emerge following issue posting and feature requests by users. Following this beta stage of development we will make the package available on CRAN for easier distribution. A critical next step will also be the development of tutorial materials to accompany the documentation, enabling new R users to quickly apply the core functions to a sampling of common use cases.

Acknowledgements

REDACTED

References